

REPCREC is a java project that models a distributed database and implements the available copies approach, strict two-phase locking, multiversion concurrency control, deadlock detection, replication and failure recovery.

Objects:

- Lock
 - Stores the ID of the transaction holding this lock, the lock start time, and whether it's a read or write lock.
- LockTable
 - Each lock table is local to a site and consists of three hashmaps (with variable IDs as keys):
 - writeLocks: each variable maps to a single transaction holding a lock
 - readLocks: each variable maps to a list of transactions holding locks
 - lockQueue: each variable maps to a list of pending locks
- Operation
 - Each line that is read in from the test file is created as an operation
 - Types: begin, beginRO, R, W, end, fail, recover, dump
- Site
 - Each site has a list of variables, an independent lockTable, and a list of updates (to keep track of commit history at this site)
 - The list of variables is initialized based on the site's number (ID: 1-10)
 - If the number is even, all even variables and 2 odd variables (number-1, number+9) are added to the site's array list of variables.
 - If the number is odd, all even variables are added to the site's array list of variables.
 - Keeps track of whether the site is currently down or has failed in the past and the time it most recently failed/recovered
 - At failure
 - the lock table is erased
 - At recovery
 - all values committed before failure time are preserved
 - non-replicated data is available for reads and writes immediately
 - replicated data is available for writes immediately, not available for reads until a write has been committed
 - getLatestValueRO(Operation o, int Time)
 - returns the latest committed value from when this operation began
- Transaction
 - An operation that reads or writes
 - Each transaction consists of a list of pending (uncommitted) operations
- TransactionManager
 - Processes each operation passed in as a list from RepCRec
 - Keeps track of 3 lists: operations, blockedOperations and sites

- Keeps track of 2 hashmaps: transactions and graph (of conflicts)
- Initializes a list of 10 sites
- simulate():
 - A while loop iterates until the lists of operations and blockedOperations are both empty
 - During each loop, time increments by 1 and an operation is either processed completely, or added to the list of blockedOperations if the read or write cannot be processed at the current time
 - At each tick, all operations in blockedOperations are processed first
 - Calls the appropriate function to process each operation based on Operation type:
 - beginTransaction(Operation o, Boolean isReadOnly)
 - adds a new Transaction object to the transactions hashmap
 - endTransaction(Operation o)
 - If the transaction is read only, do nothing.
 - Else, abort the transaction if necessary
 - Else, commit all values from the pending operations for this transaction
 - Drops all locks and clears all conflicts for this transaction
 - write(Operation o)
 - Checks if sites are down or locked depending on the variable
 - If it can successfully retrieve the locks for this transaction, calls addGraphConflicts and return true
 - Else, adds the transaction to the lock queue and returns false
 - read(Operation o)
 - If the transaction is read only, process it if possible
 - Checks if sites are down or locked depending on the variable
 - If it can successfully retrieve the locks for this transaction, calls addGraphConflicts and return true
 - Else, adds the transaction to the lock queue and returns false
 - dump(Operation o)
 - Output the state depending on whether a dumpVariable or dumpSite was provided
 - failSite(Site s)
 - Calls site.fail(currentTime)
 - recoverSite(Site s)
 - Calls site.recover(currentTime)
 - After a read or write operation is processed, calls detectDeadlock():
 - Calls detectCycleStart(graph):
 - It returns the cycle, or an empty list depending on whether or not there is a cycle in the graph
 - If there is no cycle, then there is no deadlock and it returns

- If there is a cycle, it will look at all the transactions, given by their ID in the arraylist, and determine which transaction is the youngest
- Once the youngest transaction is detected, it is aborted and its locks and conflicts are cleared
- Update
 - Stored at each site per variable
 - Stores the value and time for each commit
- Variable
 - Has an ID, value and latest commitTime

